

TensorCore

MXU
128×128

MXU
128×128

MXU
128×128

MXU
128×128

VPU (8×128 lanes)
+ VREGs

scalar unit
+ SMEM

VMEM (128 MiB on v5e)
programmer-managed scratchpad

HBM

16 GiB (v5e)
96 GB (v5p)

BW to core:
0.8 TB/s (v5e)
2.8 TB/s (v5p)

VMEM ↔ compute $\approx 22\times$ HBM BW