

VMEM ↔ MXU ($\approx 22\times$ HBM)

18.0 TB/s

HBM ↔ TensorCore

820 GB/s

ICI, all 4 links out

180 GB/s

ICI, one link one way

45 GB/s

PCIe to host

16 GB/s

DCN (per chip)

3 GB/s

10^9 10^{10} 10^{11} 10^{12} 10^{13} 10^{14}

bytes / s (TPU v5e, log scale)