

training GFLOPs per token

- MLP
- attn. projections
- LM head
- attention scores

% = attention-score share
84%

250

200

150

100

50

0

2k

4k

8k

32k

128k

sequence length s

8%

14%

25%

57%

84%

