

training GFLOPs per token

- MLP
- attn. projections
- LM head
- attention scores

% = attention-score share
84%

