

one SM: 4 subpartitions

Tensor Core
(matmul unit)

Tensor Core
(matmul unit)

Tensor Core
(matmul unit)

Tensor Core
(matmul unit)

warp sched.
32 CUDA cores

warp sched.
32 CUDA cores

warp sched.
32 CUDA cores

warp sched.
32 CUDA cores

registers
64 KiB

registers
64 KiB

registers
64 KiB

registers
64 KiB

SMEM / L1 256 kB per SM (programmer-managed)

SM

SM

SM

SM

SM

SM

... ×132

host CPU

PCIe

L2 cache 50 MB, shared by all SMs · ~5.5 TB/s

HBM (off-chip) 80 GB · 3.35 TB/s