


attention, linear $11sbh$

attention scores $5as^2b$

MLP $19sbh$


LayerNorms $4sbh$

$s = 1k$




0.31 GB

$s = 4k$



3.3 GB

$s = 16k$



45 GB

0

25%

50%

75%

100%

share of per-layer bytes (total at right)