

attention, linear 11sbh

attention scores $5as^2b$

MLP 19sbh


LayerNorms 4sbh

s = 1k




0.31 GB

s = 4k



3.3 GB

s = 16k



45 GB

0

25%

50%

75%

100%

share of per-layer bytes (total at right)