




fp32: weights, optimizer state, reductions, unscale
half: matmul inputs, saved activations, activation grads

with bf16, $S = 1$ and the scale/skip boxes drop out